

Concrete_Compressive_Strength_Prediction

December 9, 2022

1 Problem Statement

- To build a regression model to predict the concrete compressive strength based on the different features

1.1 Data Description

- Given is the variable name, variable type, the measurement unit and a brief description
- The concrete compressive strength is the regression problem. The order of this listing corresponds to the order of numerals along the rows of the database.

Data			
Name	Type	Measure	Description
Cement (com- po- nent 1)	quantitative	kg in a m3 mixture	Input Variable
Blast Fur- nace Slag (com- po- nent 2)	quantitative	kg in a m3 mixture	Input Variable
Fly Ash (com- po- nent 3)	quantitative	kg in a m3 mixture	Input Variable
Water (com- po- nent 4)	quantitative	kg in a m3 mixture	Input Variable

Data			
Name	Type	Measure	Description
Superplasticizer in concrete (component 5)	quantitative	kg in a m3 mixture	Input Variable
Coarse aggregate (component 6)	quantitative	kg in a m3 mixture	Input Variable
Fine aggregate (component 7)	quantitative	kg in a m3 mixture	Input Variable
Age	quantitative	Day (1~365)	Input Variable
Concrete compressive strength	quantitative	Mpa	Output Variable

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import scipy.stats as stats
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

2 Data Ingestion

Dataset **Link** : * <https://archive.ics.uci.edu/ml/machine-learning-databases/concrete/compressive/>

```
[3]: df = pd.read_csv('/content/drive/MyDrive/FSDS_Job_Guarantee/Projects/Cement_Strength Prediction/New Project/Concrete_Dataset.csv')
```

```
[4]: df
```

```
[4]:      Cement (component 1)(kg in a m^3 mixture)  \
0                                           540.0
1                                           540.0
2                                           332.5
3                                           332.5
4                                           198.6
...                                           ...
1025                                          276.4
1026                                          322.2
1027                                          148.5
1028                                          159.1
1029                                          260.9

      Blast Furnace Slag (component 2)(kg in a m^3 mixture)  \
0                                           0.0
1                                           0.0
2                                          142.5
3                                          142.5
4                                          132.4
...                                           ...
1025                                          116.0
1026                                           0.0
1027                                          139.4
1028                                          186.7
1029                                          100.5

      Fly Ash (component 3)(kg in a m^3 mixture)  \
0                                           0.0
1                                           0.0
2                                           0.0
3                                           0.0
4                                           0.0
...                                           ...
1025                                          90.3
1026                                          115.6
1027                                          108.6
1028                                           0.0
```

1029 78.3

	Water (component 4)(kg in a m ³ mixture) \
0	162.0
1	162.0
2	228.0
3	228.0
4	192.0
...	...
1025	179.6
1026	196.0
1027	192.7
1028	175.6
1029	200.6

	Superplasticizer (component 5)(kg in a m ³ mixture) \
0	2.5
1	2.5
2	0.0
3	0.0
4	0.0
...	...
1025	8.9
1026	10.4
1027	6.1
1028	11.3
1029	8.6

	Coarse Aggregate (component 6)(kg in a m ³ mixture) \
0	1040.0
1	1055.0
2	932.0
3	932.0
4	978.4
...	...
1025	870.1
1026	817.9
1027	892.4
1028	989.6
1029	864.5

	Fine Aggregate (component 7)(kg in a m ³ mixture)	Age (day) \
0	676.0	28
1	676.0	28
2	594.0	270
3	594.0	365
4	825.5	360

...
1025	768.3	28
1026	813.4	28
1027	780.0	28
1028	788.9	28
1029	761.5	28

Concrete compressive strength(MPa, megapascals)		
0	79.99	
1	61.89	
2	40.27	
3	41.05	
4	44.30	
...	...	
1025	44.28	
1026	31.18	
1027	23.70	
1028	32.77	
1029	32.40	

[1030 rows x 9 columns]

Rename Columns

```
[5]: df.rename(columns = {
df.columns[0] : 'Cement',
df.columns[1] : 'Blast_Furnace_Slag',
df.columns[2] : 'Fly_Ash',
df.columns[3] : 'Water',
df.columns[4] : 'Superplasticizer',
df.columns[5] : 'Coarse_Aggregate',
df.columns[6] : 'Fine_Aggregate',
df.columns[7] : 'Age(day)',
df.columns[8] : 'Concrete_Compressive_Strength'
}, inplace = True)
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Cement                                1030 non-null   float64
1   Blast_Furnace_Slag                    1030 non-null   float64
2   Fly_Ash                               1030 non-null   float64
3   Water                                  1030 non-null   float64
4   Superplasticizer                      1030 non-null   float64
```

```

5   Coarse_Aggregate          1030 non-null   float64
6   Fine_Aggregate            1030 non-null   float64
7   Age(day)                   1030 non-null   int64
8   Concrete_Compressive_Strength 1030 non-null   float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB

```

```
[7]: df.describe().T
```

```

[7]:
      count      mean      std      min  \
Cement      1030.0  281.167864  104.506364  102.00
Blast_Furnace_Slag 1030.0   73.895825   86.279342    0.00
Fly_Ash      1030.0   54.188350   63.997004    0.00
Water        1030.0  181.567282   21.354219  121.80
Superplasticizer 1030.0    6.204660    5.973841    0.00
Coarse_Aggregate 1030.0  972.918932   77.753954  801.00
Fine_Aggregate 1030.0  773.580485   80.175980  594.00
Age(day)      1030.0   45.662136   63.169912    1.00
Concrete_Compressive_Strength 1030.0   35.817961   16.705742    2.33

      25%      50%      75%      max
Cement    192.375  272.900  350.000  540.0
Blast_Furnace_Slag 0.000   22.000  142.950  359.4
Fly_Ash    0.000    0.000  118.300  200.1
Water      164.900  185.000  192.000  247.0
Superplasticizer 0.000    6.400   10.200   32.2
Coarse_Aggregate 932.000  968.000 1029.400 1145.0
Fine_Aggregate  730.950  779.500  824.000  992.6
Age(day)       7.000   28.000   56.000  365.0
Concrete_Compressive_Strength 23.710  34.445  46.135   82.6

```

```
[8]: df.isnull().sum()
```

```

[8]: Cement      0
     Blast_Furnace_Slag 0
     Fly_Ash      0
     Water        0
     Superplasticizer 0
     Coarse_Aggregate 0
     Fine_Aggregate 0
     Age(day)      0
     Concrete_Compressive_Strength 0
dtype: int64

```


```
[9]: df.duplicated().sum()
```

```
[9]: 25
```

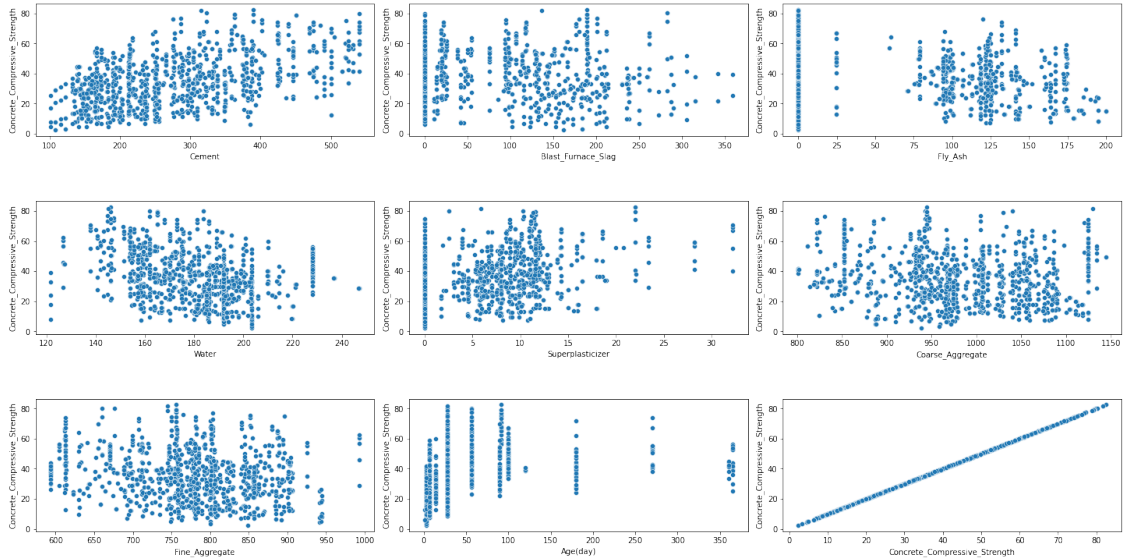
```
[10]: df.nunique()
```

```
[10]: Cement                278
      Blast_Furnace_Slag    185
      Fly_Ash              156
      Water                195
      Superplasticizer     111
      Coarse_Aggregate     284
      Fine_Aggregate       302
      Age(day)             14
      Concrete_Compressive_Strength  845
      dtype: int64
```

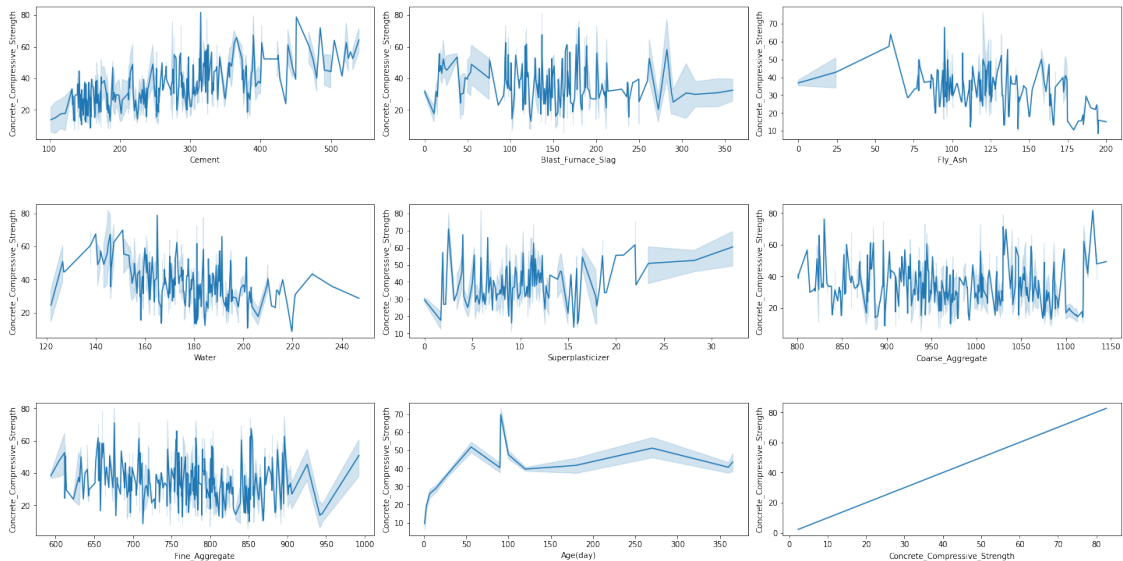
```
[11]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,5))
      index = 0
      ax = ax.flatten()
      for col, value in df.items():
          sns.distplot(value, ax=ax[index])
          index += 1
      plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[12]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,10))
      index = 0
      ax = ax.flatten()
      for col, value in df.items():
          sns.scatterplot(df[col], df['Concrete_Compressive_Strength'], ax=ax[index])
          index += 1
      plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




```
[13]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.lineplot(df[col], df['Concrete_Compressive_Strength'], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




```
[14]: plt.figure(figsize = (18,8))
sns.heatmap(df.corr().abs(), vmin = -0.5, vmax = 1, annot=True)
```


[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7fce924b8bb0>



3 Outlier Handling

```
[15]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.boxplot(y = col, data = df, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[16]: def out1(x):

    Q3 = np.nanpercentile(x , 75)
    Q1 = np.nanpercentile(x , 25)
    IQR = Q3 - Q1
    UL = (IQR * 1.5 + Q3)
    LL = (Q1 - IQR * 1.5)
    out = [1 if (a > UL) | (a < LL) else 0 for a in x]
    return(out)

# check #outliers in each variable
print(df.apply(out1).apply(sum))

# Function to Replace outlier with LL / UL

def out_impute(x):

    Q3 = np.nanpercentile(x , 75)
    Q1 = np.nanpercentile(x , 25)
    IQR = Q3 - Q1
    UL = (IQR * 1.5 + Q3)
    LL = (Q1 - IQR * 1.5)
    xnew = [x.mean() if (a<=LL) else x.mean() if (a>=UL) else a for a in x]
    return(xnew)

df = df.apply(out_impute) # Create new data with inputed values
```

```
Cement          0
Blast_Furnace_Slag  2
```

```

Fly_Ash                0
Water                  9
Superplasticizer      10
Coarse_Aggregate       0
Fine_Aggregate         5
Age(day)               59
Concrete_Compressive_Strength  4
dtype: int64

```

```
[17]: df
```

```

[17]:      Cement  Blast_Furnace_Slag  Fly_Ash  Water  Superplasticizer  \
0      540.0          0.0          0.0  162.0          2.5
1      540.0          0.0          0.0  162.0          2.5
2      332.5        142.5          0.0  228.0          0.0
3      332.5        142.5          0.0  228.0          0.0
4      198.6        132.4          0.0  192.0          0.0
...
1025    276.4          116.0         90.3  179.6          8.9
1026    322.2           0.0        115.6  196.0         10.4
1027    148.5        139.4        108.6  192.7          6.1
1028    159.1        186.7          0.0  175.6         11.3
1029    260.9        100.5         78.3  200.6          8.6

```

```

      Coarse_Aggregate  Fine_Aggregate  Age(day)  \
0          1040.0          676.0  28.000000
1          1055.0          676.0  28.000000
2           932.0          594.0  45.662136
3           932.0          594.0  45.662136
4           978.4          825.5  45.662136
...
1025          870.1          768.3  28.000000
1026          817.9          813.4  28.000000
1027          892.4          780.0  28.000000
1028          989.6          788.9  28.000000
1029          864.5          761.5  28.000000

```

```


      Concrete_Compressive_Strength
0          35.817961
1          61.890000
2          40.270000
3          41.050000
4          44.300000
...
1025          44.280000
1026          31.180000
1027          23.700000

```

```
1028                 32.770000
1029                 32.400000
```

```
[1030 rows x 9 columns]
```

```
[18]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.boxplot(y = col, data = df, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



4 Iteration Checks

```
[19]: """Can be Deleted"""
# df.drop(['Coarse_Aggregate', 'Fine_Aggregate'], axis=1, inplace=True)
```


```
[19]: 'Can be Deleted'
```

5 Transformation

Log: Log transformation helps reducing skewness when you have skewed data.

```
[20]: for col in ['Blast_Furnace_Slag', 'Fly_Ash', 'Superplasticizer']:
    mean_value = df[col].mean()
    df[col] = np.log(df[col])
    df[col] = df[col].mask(np.isinf(df[col])).fillna(mean_value)
# sns.distplot(df['Blast_Furnace_Slag'])
```

```
[21]: fig, ax = plt.subplots(ncols=3, nrows=1, figsize=(20,5))
index = 0
ax = ax.flatten()
for col in ['Blast_Furnace_Slag', 'Fly_Ash', 'Superplasticizer']:
    sns.distplot(df[col], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



6 Jarque Bera Test - Normality Check

- The Jarque-Bera Test is a test to determine if a set of data values follows the normal distribution based on the data's skewness and kurtosis
- $JB = n/2[S^2 + 1/4(K - 3)^2]$
- Where
 - n = the number of values for the data
 - S is the sample skewness (how much the data leans away from the mean)
 - K is the sample kurtosis (how thick the tails of the distribution are)
- The test statistic result will always be greater than or equal to zero since:
 1. The sample skewness in the test statistic equation is always squared, meaning S^2 is always positive or zero.
 2. The sample kurtosis is always positive or zero since the numerator is raised to the 4th power and the denominator is squared.
 3. The difference between the sample kurtosis and 3 is squared, meaning this term of the test statistic equation is always positive or zero.
 4. The sum of two terms ≥ 0 will also be greater than or equal to zero.

- We know if our data follows a normal distribution if the test statistic is close to zero and the p-value is larger than our standard 0.05. The p-value relates to a null hypothesis that the data is following a normal distribution. If the test statistic is large and the p-value is less than 0.05, the data does not follow a normal distribution.

```
[22]: from scipy.stats import jarque_bera
      for col in df.columns:
          print(f"{col} : {jarque_bera(df[col])[1]}")
```

```
Cement : 6.220579606974752e-13
Blast_Furnace_Slag : 0.0
Fly_Ash : 0.0
Water : 0.44805110686953276
Superplasticizer : 0.0
Coarse_Aggregate : 0.00036585771571151504
Fine_Aggregate : 3.5654006097618485e-05
Age(day) : 0.0
Concrete_Compressive_Strength : 1.6120992951673685e-07
```

Observation: * Except water every features has normal distribution beacause p-value is less than 0.05

7 Segregating Independent and Dependent Variable

```
[23]: X = df.iloc[ : , :-1]
      y = df.iloc[ : , -1:]
```

7.0.1 Checking Significant features based on P-values

- **H0: Column/Feature does not affect concrete strength**
- **H1: Column/Feature affects concrete strength**
- So, if a column shows p-value ≤ 0.05 then we reject the null hypothesis and say that ‘Column/Feature affects medical expenses.’
- We don’t have to actually calculate p-values for each and every column. We can simply use OLS from statsmodels.api which basically helps to fit linear regression model and also lets us know what the p-values are.

```
[25]: import statsmodels.api as sm
      X_check = sm.add_constant(X)
      model = sm.OLS(y,X_check)
      results = model.fit()
      print(results.summary())
```

OLS Regression Results

```

=====
Dep. Variable:      Concrete_Compressive_Strength    R-squared:
0.727
Model:              OLS                            Adj. R-squared:
0.725
Method:            Least Squares                   F-statistic:
339.7
Date:              Fri, 09 Dec 2022                 Prob (F-statistic):
1.46e-281
Time:              20:31:44                         Log-Likelihood:
-3677.8
No. Observations:  1030                            AIC:
7374.
Df Residuals:      1021                            BIC:
7418.
Df Model:          8
Covariance Type:   nonrobust
=====

```

	coef	std err	t	P> t	[0.025
0.975]					

const	117.3157	10.926	10.737	0.000	95.875
138.756					
Cement	0.0679	0.004	18.798	0.000	0.061
0.075					
Blast_Furnace_Slag	-0.0854	0.011	-7.675	0.000	-0.107
-0.064					
Fly_Ash	0.0480	0.017	2.779	0.006	0.014
0.082					
Water	-0.2503	0.021	-11.970	0.000	-0.291
-0.209					
Superplasticizer	-1.4645	0.237	-6.189	0.000	-1.929
-1.000					
Coarse_Aggregate	-0.0258	0.005	-5.185	0.000	-0.036
-0.016					
Fine_Aggregate	-0.0440	0.005	-8.363	0.000	-0.054
-0.034					
Age(day)	0.3153	0.010	32.663	0.000	0.296
0.334					

```

=====
Omnibus:          26.362    Durbin-Watson:          1.353
Prob(Omnibus):    0.000    Jarque-Bera (JB):      39.152
Skew:             0.242    Prob(JB):              3.15e-09
Kurtosis:         3.824    Cond. No.              5.24e+04
=====

```

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.24e+04. This might indicate that there are strong multicollinearity or other numerical problems.

8 Train Test Split

```
[26]: from sklearn.model_selection import train_test_split
```

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳ random_state=42)
```

9 Scaling

```
[28]: from sklearn.preprocessing import RobustScaler # In case of outlier
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
```

```
[29]: scaler = MinMaxScaler()
```

```
[30]: X_train_Scaled = scaler.fit_transform(X_train)
```

```
[31]: X_test_Scaled = scaler.transform(X_test)
```

10 Multicollinearly Check

```
[32]: ### VIF Check
X_train = pd.DataFrame(X_train)
from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = [variance_inflation_factor(X_train.values, i) for i in range(X_train.
↳ shape[1])]
for i,j in zip(df.columns,vif):
    print(f"{i} : {j}")
```

```
Cement : 9.236023302973546
Blast_Furnace_Slag : 2.693939762131527
Fly_Ash : 6.657528031015882
Water : 64.46210482974766
Superplasticizer : 11.52022422231611
```


Coarse_Aggregate : 87.9532452428649
Fine_Aggregate : 56.65132233337929
Age(day) : 2.468636766841692

```
[33]: """We are not removing correlated variable because we these are important_
      ↪features for our result"""
      ### Remove var with high VIF one by one in while loop
      # X_train = X_train2
      # while (max(vif) > 5):
      #     indx = vif.index(max(vif)) #Get the index of variable with highest VIF
      #     print(X_train.columns)
      #     X_train.drop(X_train.columns[indx],axis = 1, inplace = True)
      #     vif = [variance_inflation_factor(X_train.values, i) for i in_
      ↪range(X_train.shape[1])]
      # print(vif)
```

```
[33]: 'We are not removing correlated variable because we these are important features
      for our result'
```

11 Model Creation

```
[34]: from sklearn.linear_model import LinearRegression
      from sklearn.linear_model import Ridge
      from sklearn.linear_model import Lasso
      from sklearn.linear_model import ElasticNet
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.ensemble import RandomForestRegressor
      from sklearn.svm import SVR
      from sklearn.ensemble import GradientBoostingRegressor
      import xgboost as xg
      from sklearn.ensemble import AdaBoostRegressor
      from sklearn.model_selection import GridSearchCV
```

```
[35]: from sklearn.metrics import mean_squared_error
      from sklearn.metrics import mean_absolute_error
      from sklearn.metrics import r2_score
```

```
[36]: '''Hyperparameters of Ridge Regression'''
      Ridge_parameters = {
          'alpha': [0.001, 0.01, 0.1, 1, 10, 100, 1000]
      }

      '''Hyperparameters of Lasso Regression'''
      Lasso_parameters = {
          'alpha' : [0.001, 0.01, 0.1, 1, 10, 100, 1000]
```

```

    }

    '''Hyperparameters of Elastic Net Regression'''
    Elastic_Net_parameters = {
        'alpha' : [0.001, 0.01, 0.1, 1, 10, 100, 1000],
        'l1_ratio' : [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1]
    }

    '''Hyperparameters of Decision Tree Regressor'''
    DTC_parameters = {
        'max_depth': np.random.randint(5,10,5),
        'max_features': ['auto', 'sqrt', 'log2'],
        'ccp_alpha': [0.1, .01, .001],
        'criterion' :['squared_error']
    }

    '''Hyperparameters of Random Forest Regressor'''
    RF_parameters = {
        'n_estimators': [300,500,600,650,700],
        'max_depth': np.random.randint(5,10,5),
        'n_estimators': [200, 500],
        'max_features': ['auto', 'sqrt', 'log2'],
        'criterion' :['squared_error']
    }

    '''Hyperparameters of Support Vector Regressor'''
    SVR_parameters = {
        'C': [0.1, 1, 10, 100, 1000],
        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
        'kernel': ['rbf']
    }

    '''Hyperparameters of Gradient Boosted Regressor'''
    GBR_parameters = {
        'min_samples_split':range(5,20,5),
        'max_depth':range(5,10,2)
    }

    '''Hyperparameters of XG Boost Regressor'''
    XG_parameters = {
        'min_child_weight': [1, 5, 10],
        'gamma': [0.5, 1, 1.5, 2, 5],
        'subsample': [0.6, 0.8, 1.0],
        'colsample_bytree': [0.6, 0.8, 1.0],
        'max_depth': [3, 4, 5]
    }

```

```

# '''Hyperparameters of AdaBoost Classifier'''
AB_parameters = {
    'base_estimator__max_depth': [i for i in range(2,11,2)],
    'base_estimator__min_samples_leaf': [5,10],
    'n_estimators': [10,50,250,1000],
    'learning_rate': [0.01,0.1]
}

'''All Models'''
models = {
1 : LinearRegression(),
2 : Ridge(alpha=1),
3 : GridSearchCV(estimator = Ridge(alpha=1), param_grid = Ridge_parameters,
    ↳scoring='r2', verbose=1, n_jobs=-1),
4 : Lasso(alpha=1),
5 : GridSearchCV(estimator = Lasso(alpha=1), param_grid = Lasso_parameters,
    ↳scoring='r2', verbose=1, n_jobs=-1),
6 : ElasticNet(alpha=1, l1_ratio=0.5),
7 : GridSearchCV(estimator = ElasticNet(), param_grid = Elastic_Net_parameters,
    ↳scoring='r2', verbose=1, n_jobs=-1),
8 : DecisionTreeRegressor(),
9 : GridSearchCV(estimator = DecisionTreeRegressor(), param_grid =
    ↳DTC_parameters, scoring='r2', verbose=1, n_jobs=-1),
10 : RandomForestRegressor(),
11 : GridSearchCV(estimator = RandomForestRegressor(), param_grid =
    ↳RF_parameters, scoring='r2', verbose=1, n_jobs=-1),
12 : SVR(),
13 : GridSearchCV(estimator = SVR(), param_grid = SVR_parameters, scoring='r2',
    ↳verbose=1, n_jobs=-1),
14 : GradientBoostingRegressor(),
15 : GridSearchCV(estimator = GradientBoostingRegressor(), param_grid =
    ↳GBR_parameters, scoring='r2', verbose=1, n_jobs=-1),
16 : xg.XGBRegressor(objective = 'reg:linear'),
17 : GridSearchCV(estimator = xg.XGBRegressor(objective = 'reg:linear'),
    ↳param_grid = XG_parameters, scoring='r2', verbose=1, n_jobs=-1),
18 : AdaBoostRegressor(),
19 : GridSearchCV(estimator = AdaBoostRegressor(), param_grid = AB_parameters,
    ↳scoring='r2', verbose=1, n_jobs=-1)
}

```

```
[37]: map_keys = list(models.keys())
```

```

[38]: # Get model name using id from linear_model_collection
def get_model_building_technique_name(num):
    if num == 1:
        return 'Linear Regression'

```

```

if num == 2:
    return 'Ridge Regression'
if num == 3:
    return 'Ridge Regression Tuned'
if num == 4:
    return 'Lasso Regression'
if num == 5:
    return 'Lasso Regression Tuned'
if num == 6:
    return 'Elastic Net Regression'
if num == 7:
    return 'Elastic Net Regression Tuned'
if num == 8:
    return 'Decision Trees Regressor'
if num == 9 :
    return 'Decison Trees Regressor Tuned'
if num == 10:
    return 'Random Forest Regressor'
if num == 11:
    return 'Random Forest Regressor Tuned'
if num == 12:
    return 'Support Vector Regressor'
if num == 13:
    return 'Support Vector Regressor Tuned'
if num == 14:
    return 'Gradient Boosting regressor'
if num == 15:
    return 'Gradient Boosting regressor Tuned'
if num == 16:
    return 'XG Boost Regressor'
if num == 17:
    return 'XG Boost Regressor Tuned'
if num == 18:
    return 'AdaBoost Regressor '
if num == 19:
    return 'AdaBoost Regressor Tuned '
return ''

```

```

[39]: """Function to calculate all evaluation metrics"""
def evaluation(test,pred, test_set):
    R2_score = r2_score(test, pred)
    adj_R2_score = 1-((1-R2_score)*(len(test)-1)/(len(test)-test_set.
    ↪shape[1]-1))
    MSE = round(mean_squared_error(test, pred),2)
    RMSE = round(np.sqrt(MSE),2)
    MAE = round(mean_absolute_error(test, pred),2)
    return round(R2_score,2),round(adj_R2_score,2),MSE,RMSE,MAE

```

```

[40]: results = [];
for key_index in range(len(map_keys)):
    key = map_keys[key_index]
    try:
        if key in [1,2,3,4,5,6,7,12,13]:
            print(key)
            model = models[key]
            model.fit(X_train, y_train)

            y_pred = model.predict(X_test)
            y_pred_train = model.predict(X_train)

            '''Test Accuracy'''
            r2_score_test, adj_r2_score_test, MSE_test, RMSE_test, MAE_test = ↵
            ↵evaluation(y_test, y_pred, X_test)

            '''Train Accuracy'''
            r2_score_train, adj_r2_score_train, MSE_train, RMSE_train, MAE_train = ↵
            ↵evaluation(y_train, y_pred_train, X_train)

            results.append({
                'Model Name' : get_model_building_technique_name(key),
                'Trained Model' : model,
                'R2_Score_Test' : r2_score_test,
                'Adj_R2_Score_Test' : adj_r2_score_test,
                'MSE_Test' : MSE_test,
                'RMSE_Test' : RMSE_test,
                'MAE_Test' : MAE_test,
                'R2_Score_Train' : r2_score_train,
                'Adj_R2_Score_Train' : adj_r2_score_train,
                'MSE_Train' : MSE_train,
                'RMSE_Train' : RMSE_train,
                'MAE_Train' : MAE_train
            })

        if key in [8,9,10,11,14,15,16,17,18,19]:
            key = map_keys[key_index]
            model = models[key]
            print(key)
            model.fit(X_train_Scaled, y_train)
            y_pred = model.predict(X_test_Scaled)
            y_pred_train = model.predict(X_train_Scaled)

            '''Test Accuracy'''
            r2_Score_test, adj_r2_score_test, MSE_test, RMSE_test, MAE_test = ↵
            ↵evaluation(y_test, y_pred, X_test)

```

```

'''Train Accuracy'''
r2_score_train, adj_r2_score_train, MSE_train, RMSE_train, MAE_train = ↳ evaluation(y_train, y_pred_train, X_train)

results.append({
    'Model Name' : get_model_building_technique_name(key),
    'Trained Model' : model,
    'R2_Score_Test' : r2_score_test,
    'Adj_R2_Score_Test' : adj_r2_score_test,
    'MSE_Test' : MSE_test,
    'RMSE_Test' : RMSE_test,
    'MAE_Test' : MAE_test,
    'R2_Score_Train' : r2_score_train,
    'Adj_R2_Score_Train' : adj_r2_score_train,
    'MSE_Train' : MSE_train,
    'RMSE_Train' : RMSE_train,
    'MAE_Train' : MAE_train
})
except Exception as e:
    print(e)

del model

```

```

1
2
3
Fitting 5 folds for each of 7 candidates, totalling 35 fits
4
5
Fitting 5 folds for each of 7 candidates, totalling 35 fits
6
7
Fitting 5 folds for each of 77 candidates, totalling 385 fits
8
9
Fitting 5 folds for each of 45 candidates, totalling 225 fits
10
11
Fitting 5 folds for each of 30 candidates, totalling 150 fits
12
13
Fitting 5 folds for each of 25 candidates, totalling 125 fits
14
15
Fitting 5 folds for each of 9 candidates, totalling 45 fits
16
[20:33:51] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear

```

```

is now deprecated in favor of reg:squarederror.
17
Fitting 5 folds for each of 405 candidates, totalling 2025 fits
[20:35:23] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear
is now deprecated in favor of reg:squarederror.
18
19
Fitting 5 folds for each of 80 candidates, totalling 400 fits
'NoneType' object has no attribute 'set_params'

```

```
[54]: result_df = pd.DataFrame(results)
```

```
[42]: result_df_test = result_df.iloc[:, [0,2,3,4,5,6]]
result_df_test
```

```
[42]:
```

	Model Name	R2_Score_Test	Adj_R2_Score_Test	\
0	Linear Regression	0.71	0.70	
1	Ridge Regression	0.71	0.70	
2	Ridge Regression Tuned	0.71	0.70	
3	Lasso Regression	0.72	0.71	
4	Lasso Regression Tuned	0.71	0.70	
5	Elastic Net Regression	0.72	0.71	
6	Elastic Net Regression Tuned	0.71	0.71	
7	Decision Trees Regressor	0.71	0.82	
8	Decison Trees Regressor Tuned	0.71	0.81	
9	Random Forest Regressor	0.71	0.87	
10	Random Forest Regressor Tuned	0.71	0.84	
11	Support Vector Regressor	0.22	0.20	
12	Support Vector Regressor Tuned	0.85	0.85	
13	Gradient Boosting regressor	0.85	0.87	
14	Gradient Boosting regressor Tuned	0.85	0.90	
15	XG Boost Regressor	0.85	0.87	
16	XG Boost Regressor Tuned	0.85	0.90	
17	AdaBoost Regressor	0.85	0.74	

	MSE_Test	RMSE_Test	MAE_Test
0	76.11	8.72	6.74
1	76.11	8.72	6.74
2	76.04	8.72	6.74
3	73.88	8.60	6.71
4	76.08	8.72	6.74
5	73.86	8.59	6.71
6	75.64	8.70	6.73
7	47.05	6.86	4.39
8	47.54	6.89	5.06
9	32.12	5.67	3.99
10	39.94	6.32	4.74

11	206.55	14.37	11.66
12	39.27	6.27	4.29
13	32.68	5.72	4.11
14	25.35	5.03	3.45
15	33.50	5.79	4.17
16	25.15	5.01	3.44
17	66.82	8.17	6.69

```
[43]: result_df_train = result_df.iloc[:, [0,7,8,9,10,11]]
result_df_train
```

```
[43]:
```

	Model Name	R2_Score_Train	Adj_R2_Score_Train \
0	Linear Regression	0.73	0.73
1	Ridge Regression	0.73	0.73
2	Ridge Regression Tuned	0.73	0.73
3	Lasso Regression	0.73	0.72
4	Lasso Regression Tuned	0.73	0.73
5	Elastic Net Regression	0.73	0.72
6	Elastic Net Regression Tuned	0.73	0.73
7	Decision Trees Regressor	0.73	1.00
8	Decison Trees Regressor Tuned	0.73	0.93
9	Random Forest Regressor	0.73	0.98
10	Random Forest Regressor Tuned	0.73	0.94
11	Support Vector Regressor	0.22	0.21
12	Support Vector Regressor Tuned	0.94	0.94
13	Gradient Boosting regressor	0.94	0.93
14	Gradient Boosting regressor Tuned	0.94	0.97
15	XG Boost Regressor	0.94	0.93
16	XG Boost Regressor Tuned	0.94	0.97
17	AdaBoost Regressor	0.94	0.78

	MSE_Train	RMSE_Train	MAE_Train
0	74.12	8.61	6.55
1	74.12	8.61	6.55
2	74.12	8.61	6.55
3	74.87	8.65	6.61
4	74.12	8.61	6.55
5	74.95	8.66	6.61
6	74.15	8.61	6.55
7	1.07	1.03	0.14
8	19.78	4.45	3.31
9	5.67	2.38	1.52
10	16.84	4.10	3.06
11	212.70	14.58	11.71
12	17.24	4.15	1.99
13	17.95	4.24	3.07
14	6.90	2.63	1.71

15	19.37	4.40	3.15
16	9.08	3.01	1.99
17	60.27	7.76	6.43

12 Save and Load Best Model

```
[44]: Best_Model_Name = result_df['Trained_
↳Model'][result_df[result_df['Adj_R2_Score_Test'] ==_
↳max(result_df['Adj_R2_Score_Test'])]['Trained Model'].index[0]]
Best_Model_Index = result_df['Trained_
↳Model'][result_df[result_df['Adj_R2_Score_Test'] ==_
↳max(result_df['Adj_R2_Score_Test'])]['Trained Model'].index].index[0]
import pickle
""" Save the model to disk """
filename = 'Concrete_Compressive_Strength_Prediction.sav'
pickle.dump(Best_Model_Name, open(filename, 'wb'))
"""Load the model from disk"""
loaded_model = pickle.load(open(filename, 'rb'))
loaded_model.fit(X_train, y_train)
y_pred = loaded_model.predict(X_test)
```

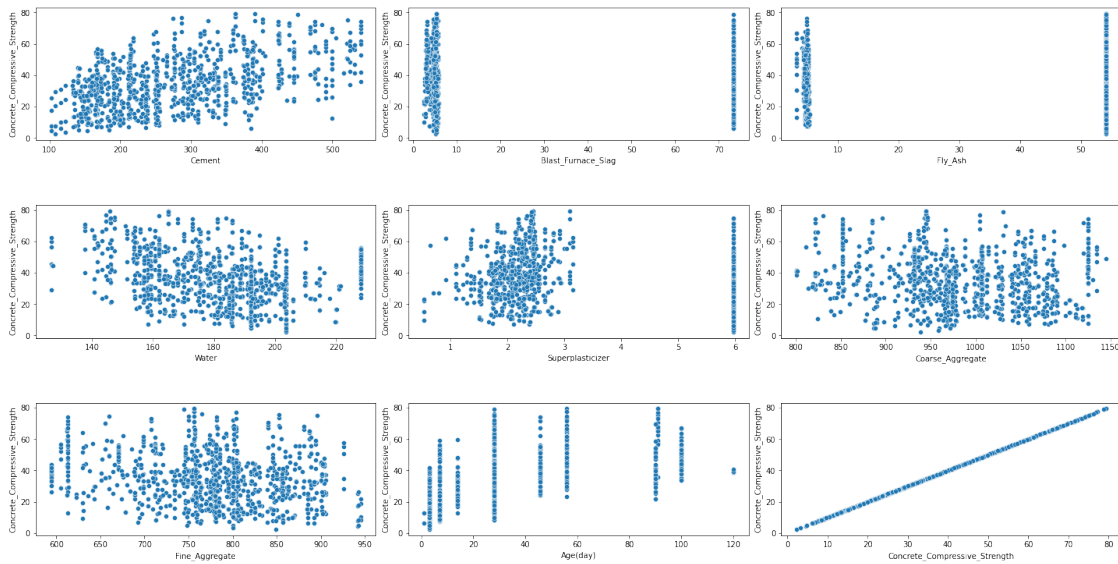
Fitting 5 folds for each of 9 candidates, totalling 45 fits

13 Linear Regression Assumptions

```
[45]: from sklearn.linear_model import LinearRegression
LR_Model = LinearRegression()
LR_Model.fit(X_train, y_train)
y_pred = LR_Model.predict(X_test)
```

13.1 1. Linear Relationship

```
[46]: fig, ax = plt.subplots(ncols=3, nrows=3, figsize=(20,10))
index = 0
ax = ax.flatten()
for col, value in df.items():
    sns.scatterplot(df[col], df['Concrete_Compressive_Strength'], ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




13.2 2. Normality


- Residuals should be normally distributed

```
[47]: residuals = y_test - y_pred
      sns.displot(residuals, kind='kde')
```

```
[47]: <seaborn.axisgrid.FacetGrid at 0x7fce96d1eaf0>
```



```
[48]: # QQ Plot
import scipy as sp
fig, ax = plt.subplots(figsize=(6,4))
sp.stats.probplot(residuals['Concrete_Compressive_Strength'], plot=ax, fit=True)
plt.show()
```



13.3 3. Multicollinearity

```
[49]: ### VIF Check
X_train = pd.DataFrame(X_train)
from statsmodels.stats.outliers_influence import variance_inflation_factor
vif = [variance_inflation_factor(X_train.values, i) for i in range(X_train.
    ↳shape[1])]
for i,j in zip(df.columns,vif):
    print(f"{i} : {j}")
```


```
Cement : 9.236023302973546
Blast_Furnace_Slag : 2.693939762131527
Fly_Ash : 6.657528031015882
Water : 64.46210482974766
Superplasticizer : 11.52022422231611
Coarse_Aggregate : 87.9532452428649
Fine_Aggregate : 56.65132233337929
Age(day) : 2.468636766841692
```

Observation: * There is multicollinearity in few variables but we decided not to remove because these are important for the concrete compressive strength prediction

13.4 4. Homoscedasticity

```
[50]: plt.scatter(y_pred, residuals)
```


```
[50]: <matplotlib.collections.PathCollection at 0x7fce96dd58b0>
```



13.5 5. Autocorrelation

```
[51]: plt.plot(residuals)
```

```
[51]: [<matplotlib.lines.Line2D at 0x7fce96d99bb0>]
```



13.5.1 Durbin Watson Test

- A test developed by statisticians professor James Durbin and Geoffrey Stuart Watson is used to detect autocorrelation in residuals from the Regression analysis
- **Autocorrelation:**
- Autocorrelation represents the degree of similarity between a given time series and a lagged version of itself over successive time intervals. Autocorrelation measures the relationship between a variable's current value and its past values
- **Assumptions of Durbin-Watson d Test**
- The errors are normally distributed with a mean value of 0
- The errors are stationary
- Null and Alternate Hypothesis of Durbin-Watson d Test
- Null Hypothesis: First order autocorrelation does not exist
- Alternate Hypothesis: First order autocorrelation exists
- The test statistic is approximately equal to $2*(1-r)$ where r is the sample autocorrelation of the residuals. Thus, the test statistic will always be between 0 and 4 with the following interpretation:
- A test statistic of 2 indicates no serial correlation.
- The closer the test statistics is to 0, the more evidence of positive serial correlation.

- The closer the test statistics is to 4, the more evidence of negative serial correlation.

```
[52]: from statsmodels.stats.stattools import durbin_watson
      gfg = durbin_watson(residuals)
      print(gfg)
```

[2.04853471]

Observation: * There is very little or no auto correlation present between the residuals

14 THE END